

Hard questions, written in the interviewer's voice

A walkthrough of the case-study interview as if Samay Jhunhunwalla is on the other side of the table — operator-sharp, marketplace-pragmatic, allergic to clever math that can't survive one booker objection. Each question has a planted follow-up and an explicit "what they're testing" so the prep is actually useful, not just intimidating.

THE INTERVIEWER, BRIEFLY

Samay leads product at Greenroom and writes Creating Value — a Substack about marketplace product judgment. His lens is operator-first: every framework has to survive contact with a real booker on a real Tuesday at 11pm. He reads case studies for **reasoning** over polish.

WHAT HE'LL PROBE FOR

- **Diagnosis discipline** — did you frame the right problem, or jump to a fancy solution?
- **Sample-size honesty** — do you know when your data is too thin to act on?
- **Booker trust** — does the UX survive the moment Mariana disagrees with the suggestion?
- **Failure-mode plan** — when the model is wrong, who eats it and how do you know?
- **Roll-forward judgment** — what's Phase 4, and why isn't it Phase 1?

Format note: Each block is one likely interviewer question plus one or two follow-ups. The **WHAT THEY'RE TESTING** line is the meta — what the question is actually probing under the surface. Use it to calibrate the answer's altitude (data vs judgment vs operational sense).

Contents

1. Diagnosis & framing

2. Data slicing & methodology

3. SGP waterfall — step by step

4. Smart Switch UX & booker trust

5. Improve-Deal & the caps argument

6. Insights pipeline & LLM judgment

7. Switch-Savings backtest

8. Insurance pricing & platform incentives

9. Failure modes & adverse cases

10. Rollout, measurement, Phase 4

01 Diagnosis & framing

Samay opens here. If the diagnosis is wrong, the rest of the memo doesn't matter — he'll spend 15 minutes here before touching the math.

Q1.1 · OPENING

Walk me through how you decided "deal type" was the bottleneck and not, say, settlement-night data entry or post-show payment latency. What did you rule out, and on what evidence?

FOLLOW-UP

If I gave you the same dataset but the wizard could already model door deals, would your diagnosis change? What would be the new bottleneck?

TESTING WHETHER YOU TRIAGED THE UNIVERSE OF FAILURE MODES BEFORE PICKING ONE. DID "WIZARD CAN'T MODEL 38% OF DEALS" WIN ON EVIDENCE OR ON CONVENIENCE?

Q1.2 · THE 38% NUMBER

You quote "38% of deals can't be settled in-app" repeatedly. Walk me through that arithmetic. Is it 38% of deal count, or 38% of gross, or 38% of *booker-hours*? Which one matters?

FOLLOW-UP

If 38% by count is actually 12% by gross, does your Phase-1 ranking still hold?

TESTING WHETHER YOU UNDERSTAND THAT COUNT, DOLLAR VOLUME, AND TIME SPENT ARE THREE DIFFERENT STORIES AND YOU PICKED THE RIGHT ONE FOR THE AUDIENCE.

Q1.3 · THE UNFLATTERING FRAME

A skeptic could read this memo and say "you built a calculator that recommends the deal type the wizard *happens* to support — that's tool-led, not booker-led." What's your response?

FOLLOW-UP

What's the case where the right answer is "leave the deal as %-of-net and improve the wizard instead"?

TESTING WHETHER YOU CAN ARGUE AGAINST YOUR OWN MEMO. HE WANTS TO SEE IF YOU'VE STRESS-TESTED THE FRAMING.

02 Data slicing & methodology

Samay's an operator — he'll pick a bucket and ask why that one, why those edges, and whether the sample is real.

Q2.1 • BUCKET EDGES

Your size buckets are \$0-1K • \$1-5K • \$5-15K • \$15K+ . Why those edges? They're round numbers but the industry's economics aren't linear. What does your data look like at \$4,800 vs \$5,200?

FOLLOW-UP

If I move the \$5K edge to \$6K, does any of the cross-tab cell logic flip?

TESTING WHETHER YOUR BUCKET EDGES ARE DEFENSIVE (ROUND NUMBERS BECAUSE THEY'RE READABLE) OR EMPIRICAL (THEY SIT ON NATURAL BREAKS IN THE HISTOGRAM).

Q2.2 • THE 24-MONTH WINDOW

Why 24 months? COVID is ~24 months back depending on when you count. The Crescent's deal mix in 2024 vs 2025 — same shape, or different?

FOLLOW-UP

If I gave you only the last 12 months, do your SGP tier counts collapse? Where?

TESTING WHETHER YOU'VE CONSIDERED STRUCTURAL BREAKS IN YOUR TRAINING DATA. TWENTY-FOUR MONTHS IS A CHOICE, NOT A DEFAULT.

Q2.3 • THE \$1-5K VS/PN CELL

Your supporting doc says the \$1-5K vs/pn cell has n=141 comparables. How many are *distinct agents*? How many distinct artists? If three agents dominate that cell, is your SGP really learning "venue economics" or "Bob from CAA's negotiating style"?

FOLLOW-UP

What's the cell concentration ratio that would make you fall back to genre instead of agent?

TESTING SAMPLE-SIZE SOPHISTICATION. SAME N=141 CAN BE 50 UNIQUE AGENTS OR 8 — COMPLETELY DIFFERENT SIGNAL QUALITY.

Q2.4 • SURVIVORSHIP BIAS

Your dataset is settled shows. What about deals that were proposed and never booked because the booker walked away? Those are missing from the training set — but they're exactly the deals SGP is trying to advise on.

FOLLOW-UP

How would you instrument Greenroom to start capturing the dead-deal data?

TESTING WHETHER YOU SEE THE SELECTION BIAS IN "WE TRAINED ON THE DEALS WE AGREED TO." THIS IS THE DEEPEST CUT AND PROBABLY THE QUESTION HE MOST WANTS TO HEAR A GOOD ANSWER TO.

03 SGP waterfall — step by step

Each step in the 7-step waterfall is debatable. He'll pick two or three.

Q3.1 · THE N=8 THRESHOLD

Step 2 fires when the cell has ≥ 8 comparables. Why 8? Why not 5, why not 20?

FOLLOW-UP

What's the confidence interval on a cell mean from $n=8$? Would you show that band to Mariana?

TESTING WHETHER YOU PICKED THE THRESHOLD EMPIRICALLY (BOOTSTRAPPED CI WIDTH CROSSES SOME LINE) OR BY FEEL.

Q3.2 · THE AGENT-HISTORY STEP

Step 3 weights by the agent's prior settlements at this venue. An agent who has done 3 cheap deals here looks "experienced" — should they?

FOLLOW-UP

What if the agent has 3 priors but two were disputed? Does that agent get more or less weight than a fresh agent?

TESTING WHETHER "AGENT FAMILIARITY" IS THE RIGHT SIGNAL OR "AGENT RELIABILITY" IS. DIFFERENT FORMULAS.

Q3.3 · THE WINNER STEP

Step 6 takes `max(guarantee, percentagePayout)`. That's pessimistic for the venue — you always pay the larger number. Why isn't this `min`, or a probability-weighted blend?

FOLLOW-UP

Show me the case where `max` systematically over-pays an artist who would have settled for less.

TESTING WHETHER YOU UNDERSTAND THE ASYMMETRIC LOSS FUNCTION — THE VENUE UNDERPAYS ONCE AND LOSES A RELATIONSHIP; OVER-PAYS AND LOSES A MARGIN POINT.

Q3.4 · THE ROUNDING CAVEAT

You document that `roundTo50` can land \$24 below the contract guarantee. That's a contract breach. Why is this in the memo instead of fixed in the code before the memo went out?

FOLLOW-UP

What's your process for catching this class of bug before it shows up to a booker?

TESTING HONESTY DISCIPLINE. HE'S NOT TESTING WHETHER YOU'LL FIX IT — HE'S TESTING WHETHER YOU'LL DEFEND THE CHOICE TO DISCLOSE IT. (HINT: YES, ALWAYS DISCLOSE.)

04 Smart Switch UX & booker trust

Q4.1 • MARIANA AT 11PM

It's settlement night. The Smart Switch suggestion says "convert this %-of-net to a \$4,200 flat." Mariana disagrees — she thinks \$3,800. Walk me through every UI affordance she has to override, and what data she sees behind that \$4,200.

FOLLOW-UP

If she overrides three times in a row on similar deals, what does the system do?

TESTING WHETHER THE SUGGESTION IS A DECISION OR A DRAFT. BOOKERS DON'T ACCEPT DECISIONS; THEY ACCEPT DRAFTS THEY CAN EDIT.

Q4.2 • THE PROVENANCE ENUM

You badge each suggestion with `sgp_engine` / `guarantee_amount` / `cell_mean` / etc. Real bookers — do they look at that badge? Or is it engineering hygiene wearing a UI hat?

FOLLOW-UP

If only 5% of bookers click it, do you keep it? At what cost?

TESTING WHETHER YOU CAN DEFEND TRANSPARENCY UX AGAINST THE "USERS DON'T CARE" OBJECTION. (SPOILER: IT'S NOT FOR THE MODAL USER, IT'S FOR THE MOMENT OF DISPUTE.)

Q4.3 • THE DOOR-HYBRID FLOOR

Door hybrid has a fixed \$500 artist floor. Why a constant and not a venue-specific calibration? The Crescent is 250-cap; would Bowery Ballroom take the same \$500?

FOLLOW-UP

What's the right input that should set the floor — capacity, average gross, or something else?

TESTING WHETHER YOU'VE THOUGHT ABOUT PHASE 4 — MULTI-VENUE. THE CURRENT PROTOTYPE IS ONE VENUE; A CONSTANT IS FINE. THE PRODUCT ISN'T.

05 Improve-Deal & the caps argument

Q5.1 • WHY CAPS-ONLY

You explicitly scoped Improve-Deal to caps-only. The redesign-v2 doc cites an April audit that said scaling caps by bucket isn't data-safe. Walk me through that audit — what was the test, what was the threshold?

FOLLOW-UP

If the audit had said caps do scale by gross, would you still draw the Smart-Switch / Improve-Deal boundary in the same place?

TESTING WHETHER THE SCOPE BOUNDARY IS PRINCIPLED OR PRAGMATIC. BOTH ARE FINE, BUT YOU HAVE TO KNOW WHICH.

Q5.2 · P75 AS A DEFAULT

Cap defaults are P75 of historical flats. P75 is generous to the venue. Why not P50? Why not P90?

FOLLOW-UP

If you imagine the agent on the other side of the table reading the same memo, where would they push back on P75?

TESTING TWO-SIDED MARKETPLACE THINKING. THE VENUE'S "REASONABLE CAP" IS THE AGENT'S "TIGHT CAP."

06 Insights pipeline & LLM judgment

Q6.1 · TWO-STAGE PIPELINE

Per-settlement enrichment, then per-cell clustering. Why two stages and not one? You're paying for two LLM calls per insight.

FOLLOW-UP

What's the cell with the fewest enriched rows that still produces a "top 5 themes" output? Is the output meaningful at n=3?

TESTING WHETHER YOU CAN ARTICULATE THE ENGINEERING TRADEOFF (PER-ROW REUSABILITY, IDEMPOTENCY) WITHOUT HIDING BEHIND IT.

Q6.2 · HALLUCINATION AT THE CELL LEVEL

A cell with three flagged deals produces a top-5 cluster. By construction at least two themes are made up. How do you keep that from looking like a confident lie?

FOLLOW-UP

What's the UX rule — minimum n before you render the themes? Do you render zero, or do you render with a confidence chip?

TESTING WHETHER YOU'VE THOUGHT ABOUT LLM OUTPUT AS A UI SURFACE, NOT A DATA ARTIFACT. THIS IS WHERE MOST CASE STUDIES FAIL — THEY TRUST THE LLM TOO MUCH.

Q6.3 · THE RUNTIME PROVIDER SWAP

The Settings tab swaps providers globally and invalidates the Insights cache. Sounds clean. What's the worst-case operational story when the booker swaps from Claude to GPT-4o mid-week?

FOLLOW-UP

How do you keep historical insights stable when the model behind them changes?

TESTING WHETHER YOU UNDERSTAND THAT "STATELESS LLM CALL" IS A LIE AT THE PRODUCT LEVEL — EVERY INSIGHT HAS A PROVENANCE.

07 Switch-Savings backtest

Q7.1 • COUNTERFACTUAL FALLACY

Switch-Savings says "if you had used Smart Switch, you would have saved \$X." But if Smart Switch had existed, the agent might have negotiated differently. How do you defend the counterfactual?

FOLLOW-UP

What's the discount you'd apply to the headline savings number to account for agent reaction?

TESTING WHETHER YOU UNDERSTAND THAT BACKTESTS IN TWO-SIDED MARKETS ARE NOT SINGLE-SIDED SIMULATIONS.

Q7.2 • THE TIME-SAVED ESTIMATE

"Time saved" is estimated from notes/signoff/recoup activity. Walk me through the formula. Is it minutes-per-recoup-thread? Per word in the signoff? How robust is that to a booker who writes long notes by habit?

FOLLOW-UP

What's the validation — did you ever have a booker say "yes, that was about 40 minutes"?

TESTING WHETHER YOUR TIME NUMBER IS A CALIBRATED ESTIMATE OR A PLAUSIBLE-LOOKING PLACEHOLDER.

08 Insurance pricing & platform incentives

Q8.1 • THE 4.32%

Product 1 prices at $4.32\% \times \$345 = \14.84 per show. Walk me through the 4.32%. Is that loss frequency, claim rate, or expected payout ratio?

FOLLOW-UP

What's the sensitivity? If 4.32% drifts to 6%, does the product still pencil?

TESTING WHETHER YOU UNDERSTAND INSURANCE MATH AT FIRST-PRINCIPLES OR YOU'RE QUOTING A NUMBER FROM A SPREADSHEET.

Q8.2 • ADVERSE SELECTION

Bookers who opt into Product 2 are the ones who think they're going to claim. That's adverse selection. How do you price it, and how do you keep healthy bookers from paying for sick ones?

FOLLOW-UP

Is there a version of this product that's mandatory at the platform level instead of opt-in? What does that mean for the pricing?

TESTING WHETHER YOU'VE ACTUALLY THOUGHT LIKE AN UNDERWRITER, NOT JUST A PM WITH AN EXCEL SHEET.

Q8.3 • THE TWO-SIDED DISCOUNT MATH

Your brief includes two-sided discounts for venues that also bring agent volume. Why is that the right structure vs flat pricing? What does it cost the platform when you mis-tier?

FOLLOW-UP

If Greenroom were a marketplace where agents pick the platform, would your tiering survive contact with the agent side?

TESTING TWO-SIDED THINKING, AGAIN — THE SAMAY SPECIALTY.

09 Failure modes & adverse cases

Q9.1 • THE BAD SUGGESTION

Smart Switch suggests a flat that's \$800 below what the agent has already verbally agreed to. Mariana accepts the suggestion. Three months later, the agent's annoyed. Who eats the \$800? Whose job is it to catch this?

FOLLOW-UP

What's the leading indicator in the data that this is happening at scale?

TESTING OPERATIONAL OWNERSHIP. PRODUCT CAN'T SHIP "ACCEPT AND WE'LL SEE" — THERE HAS TO BE A FEEDBACK LOOP.

Q9.2 • THE GAMED AGENT

An agent figures out the SGP formula and packs their early deals at the venue to inflate their "history weight." Six months in, SGP suggests \$5,200 because that agent has trained the algorithm. How do you detect it?

FOLLOW-UP

What's the right monitoring metric — agent-level dispersion of suggested-vs-actual?

TESTING ADVERSARIAL THINKING. DO YOU ASSUME AGENTS ARE HONEST, OR DO YOU ASSUME AGENTS ARE OPTIMIZERS?

Q9.3 • THE EMPTY-CELL CASE

A brand-new venue onboards. Day 1, zero history. What does Smart Switch say?

FOLLOW-UP

How long is the "cold start" phase, and what's the product surface during it?

TESTING WHETHER YOUR PRODUCT IS A STATIC ALGORITHM OR A SYSTEM THAT GROWS UP.

10 Rollout, measurement, Phase 4

Q10.1 • THE ROLLOUT ORDER

Your three phases are Smart Switch → Improve-Deal → Insurance. Defend the order. Why not insurance first as a wedge product?

FOLLOW-UP

What's the bookers-have-trust prerequisite that makes insurance Phase 3 instead of Phase 1?

TESTING WHETHER THE ORDER IS PRODUCT-LOGICAL (EACH PHASE EARNS THE TRUST THE NEXT ONE SPENDS) OR JUST CONVENIENT.

Q10.2 • THE SUCCESS METRIC

In six months, how do I know Phase 1 worked? Give me one leading and one lagging metric — and tell me the threshold for "kill it."

FOLLOW-UP

If acceptance rate is 70% but override rate on accepted suggestions is 40%, is that a win or a loss?

TESTING WHETHER YOUR KPI HAS A TEETH — KILL-THRESHOLD, NOT JUST A GREEN CHECK.

Q10.3 • PHASE 4

You don't write Phase 4 in the memo. Tell me what it is and why it isn't Phase 1.

FOLLOW-UP

What's the data you don't have yet that Phase 4 needs?

TESTING WHETHER YOU HAVE A REAL ROADMAP OR THREE PHASES ENDING IN FOG. THIS IS THE CLOSE — LEAVE WITH A VISION.

Q10.4 • THE EXISTENTIAL ONE

If Greenroom never builds any of this — Smart Switch, SGP, insurance — what's the actual cost to The Crescent over a year? Convince me this matters more than the next three things on our roadmap.

FOLLOW-UP

If I told you the roadmap also includes "agent-side messaging," which one wins and why?

TESTING WHETHER YOU CAN SELL THE WORK. HE RUNS PRODUCT — HE DOESN'T JUST WANT IT TO BE SMART, HE WANTS IT TO BE THE RIGHT THING TO BUILD NEXT.

How to use this in the room

- **Open with the diagnosis answer.** If Q1.1 lands, the next 30 minutes are easier. Have the 38%-by-count vs 38%-by-gross numbers in your head before you walk in.
- **Concede early on the rounding caveat (Q3.4).** Don't defend it. He's testing honesty, not engineering. The right answer is "you're right, that's a contract-breach risk, here's the one-line fix and here's why I disclosed it rather than hide it."
- **Don't bluff the insurance math (Q8).** If you don't remember the 4.32% derivation cold, say so and pull it up — the supporting doc has it. Confidence about the wrong number is the worst signal you can send to him.
- **For Q2.4 (survivorship bias)** — name the bias by name. He'll relax visibly.
- **For Q6.2 (LLM cluster honesty)** — say "minimum-n threshold" out loud. That's the answer.
- **Close on Phase 4 (Q10.3).** Have a real one. "Multi-venue SGP calibration" + "agent-side fairness audit" is a defensible pair.

Greenroom · Interview Prep · May 2026 · Built against the live prototype at commit `30ca767` + the final memo + the v2.1 redesign + the supporting + the SGP-vs-flat + the insurance pricing brief in `exports/`.